



# Advanced Distributed Systems

---

Author: Jose 胡冠洲 @ ShanghaiTech

Teacher: [Prof. Tao Xie](#) from San Diego State

Figures come from professor's slides, unless otherwise specified.

## Advanced Distributed Systems

- Distributed Systems Basics

  - Characteristics of Distributed Systems

  - Challenges & Issues

- System Models

  - What is a Model

  - Architectural Models

  - Fundamental Models

- Interprocess Communication (IPC)

  - Middleware Layer of Communication

  - Request-Reply Protocol

  - RMI & RPC

  - Events & Notifications

- Distributed File Systems

  - FS & DB Terms

  - File Service Design

- Peer-to-Peer Systems (P2P)

  - Overview of P2P Systems

  - Overlay Networks

- Time & Global State

  - Distributed Timing

  - Global State

- Coordination & Agreement

  - Distributed Mutual Exclusion

  - Election Algorithms

  - Multicast Algorithms

  - Consensus Algorithms

- Transactions Interface

  - Definition of Transactions

  - Life History of a Transaction

  - Concurrency Control

  - Recoverability from Aborts

  - Nested Transactions

- Replications

  - Replication Transparency

  - Fault-Tolerant Services

  - Highly Available Services

- Security

## Distributed Systems Basics


---

## Characteristics of Distributed Systems

Definition: **Independent** components (may be far apart) working **collaboratively** by passing messages over computer networks.

1. Concurrency
2. Lack of a global clock
3. Independent failures

Compared to Parallel Computing (focusing on multiple homogeneous processors in ONE computer), distributed systems require a better understanding of: *Heterogeneity, Openness, Security, Scalability, Failure handling, Concurrency, Transparency*.



Coulouris, Dollimore and Kindberg Distributed Systems: Concepts and Design Edn. 4  
© Pearson Education 2005

## Challenges & Issues

- Heterogeneity: Different hardware, OS, network architecture, languages; VMs
- Openness: Standard published interfaces; RFC
- Security: Confidentiality + Integrity + Availability
- Scalability: Extensiveness, cost v.s. performance
- Fault Tolerance: Detecting faults, masking (hiding), ignoring, recovery; Dependence via Redundancy
- Concurrency issues
- Transparency:
  - Access Transparency
  - Location Transparency
  - Concurrency Transparency
  - Replication Transparency
  - Failure Transparency
  - Mobility Transparency
  - Performance Transparency
  - Scaling Transparency


# System Models

"Distributed Systems: Concepts & Design", Chapter 2


## What is a Model

A description of a complex entity or process, simplified by ignoring certain details.

- *Architectural* models: focusing on distribution & communication of data / tasks amongst physical nodes



- *Fundamental* models: focusing on conceptual issues to be solved



## Architectural Models

Follow 3 steps to build an architectural model:

1. Functions of individual components
2. Placement of the components across a network
3. Interrelationships between the components

A distributed system usually expose to users **services**, which masks over heterogeneity of the underlying platform. **Servers** (a process that accepts requests from other processes) provide services, and **clients** access services.

- *Clients-Server* pattern - simple
  - 1 / more servers
  - 1 / multi-level servers; caches
  - Run code remotely v.s. Retrive code and run locally
- *Peer-to-Peer* (P2P) pattern - decentralization, improves scalability

Design requirement of distributed architectures:

- Performance: throughput, load balancing
- Quality of Service (**QoS**): latency
- Dependability: safety

# Fundamental Models

Fundamental models describe basic issues (shared among different architectures).

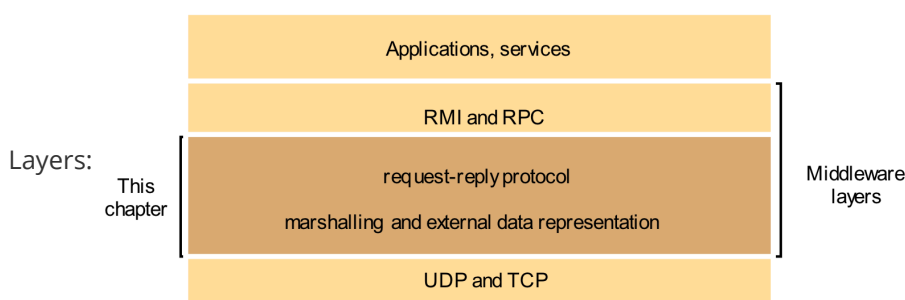
- Interaction model
  - Sequential v.s. Distributed timing / state
  - *Synchronous* v.s. *Asynchronous*
- Failure model
  - **Crashes** / Fail-stops
  - **Omission Failures**: failed to do what it is supposed to do (end up doing nothing); can recover by reapply
  - **Timing failures**: takes longer than time bound
  - **Arbitrary Failures** (*Byzantine*): inconsistent failures at arbitrary times
- Security model
  - Authorization: access rights
  - Authentication: identity
  - Denial of Service (**DoS**) attacks → Distributed DoS (**DDoS**)

## Interprocess Communication (IPC)

---

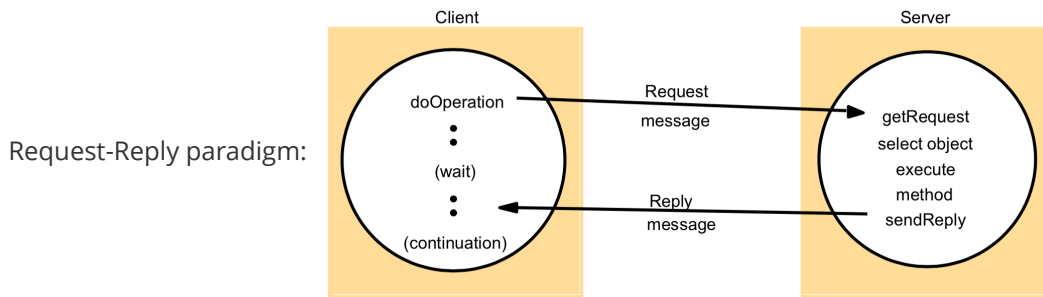
"Distributed Systems: Concepts & Design", Chapter 4

### Middleware Layer of Communication



- IPC relies on Remote Method Invocations (**RMI**) / Remote Procedure Calls (**RPC**) / Events
- RMI / RPC / Events are implemented upon **Request-Reply** protocol
- Request-Reply protocol is built upon User Datagram Protocol (UDP) / Transmission Control Protocol (TCP, *streams*)
  - Defines external data representation, e.g. HTTP, XDR, Java object serialization
- UDP / TCP are provided by the operating system, they must act on proper hardware, with a message destination
  - Destination can be IP address + *Port* / multicast
  - Often utilized through the **Socket** programming API

### Request-Reply Protocol




Involves *send* & *receive* operations, and can be synchronous / asynchronous. Some techniques can be used to protect / optimize this procedure:

- *Idempotent* operations: can be performed repeatedly with the same effect as performing once
- Cache history of replies to avoid re-execution on the same requests
- Timeout & Retrying
- ...

*Marshalling & Unmarshalling*: convert internal data to / from standard external transferring format.

## RMI & RPC

Local v.s. Remote invocation pattern:



RMI needs:

- *Message structure*: defines how a message is composed  
Example: | Message type | Request ID | Remote object reference |
- *Remote object reference*: an identifier throughout the distributed system to access the remote object  
Example: | IP address | Port # | Time | Object ID | Remote interface |
- *Remote interface*: specifying which methods can be invoked remotely, and their signatures  
Example: implemented using Interface Definition Languages (*IDL*)

Three possible effects of RMI:

1. The state of the receiver may get changed
2. A new object may be instantiated
3. Further chained invocations of other methods

Three fault-tolerance semantics choices:

Retransmit Request Message on Failure?	Duplicate Request Filtering?	Re-execute / Retransmit Reply on Duplication?	Invocation Semantics
No	\	\	<i>Maybe</i>
Yes	No	Re-execute the procedure	<i>At-least-once</i>
Yes	Yes	Retransmit reply	<i>At-most-once</i>


- *Maybe* is suitable when occasional failed invocations are acceptable
- *At-least-once* is suitable for idempotent invocations
- *At-most-once* is completely fault tolerant

A **Proxy** on the client side plays the role of a local object gate as to the invoking object, it marshalls the request and sends it downwards & unmarshalls the received result. (In RPCs, called a **Stub**)

A **Binder** is a table that maps remote object textual names to remote object references / maps remote procedures to local ports. Servers register into the binder, and clients lookup this table.

To void polling, often uses **Callbacks** (Server calls a remote call-back method on client call-back object on reply).

## Events & Notifications



Components of this pattern:

- *Object of Interest*: the object that experiences changes of state
- *Event*: occurs at the end of the state change of an Object of Interest
- *Notification*: a marshalled message containing information about an Event
- *Subscriber*: an object that is interested in a certain type of Events; will receive Notifications of that type of Events
- *Publisher*: the object responsible for generating Notifications
  - Can be the Object of Interest, or
  - An *Observer Object (Proxy)*: an object gate, like proxy in RMI, but at the server side

## Distributed File Systems

"Distributed Systems: Concepts & Design", Chapter 12

### FS & DB Terms

**Filesystem (FS):**

- Local FS [Kernel-level]: part of the operating system kernel, which wraps over naked storage devices (HDD, SSD, ...) and provides POSIX API for upper applications to use these devices; e.g. NTFS (Win), Ext4, XFS, BtrFS, ... (Linux), HFS, APFS (Apple)
- Virtual FS (VFS) [Kernel-level]: an OS module that supports interactions with multiple devices formatted as different FS formats; e.g. Linux VFS supports reading and writing over Ext4 and also a plugged-in FAT-32 USB drive (may sometimes be considered as part of the local FS)
- FS Middleware [User-level]: a middle layer between user applications and the local FS, which often provides parallel / distributed optimizations for IO scheduling to improve FS performance; e.g. GPFS, Lustre, Sun NFS, OrangeFS (PVFS), Ceph, ...
  - Can be integrated into the OS kernel as well, like NFS protocol integrated in Linux kernels
  - Can run as FUSE (Filesystem in User Space)
  - Can be implemented as an application-level dynamic library; may require user applications to get re-compiled with the designed APIs instead of POSIX

#### Database (DB):

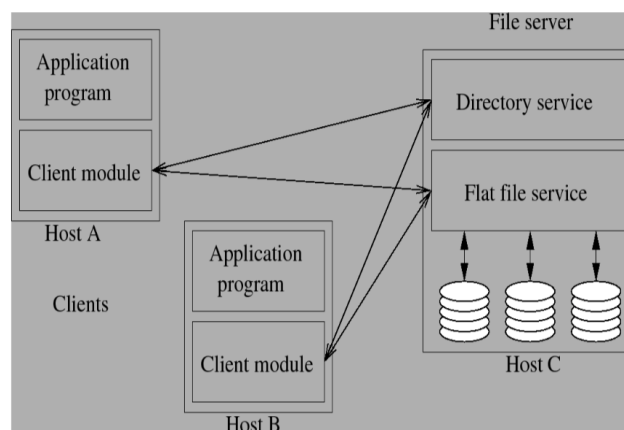
- Application-level DB [User-level]: similar to FS Middleware, provides a new layer of data management over the local FS to unify data format and provide in-memory caching; e.g. MySQL
- Naked-device DB [Kernel-level]: databases that directly operate on naked devices to improve performance, started supported by Oracle (no longer popular)
- Object Storage (Key-Value Store, KV Store) [Kernel-level / User-level]: databases built upon new object storage devices (or somehow simulated using traditional disks), guided by *NoSQL* design; e.g. Redis, Oracle NoSQL

Distributed FS we discuss here are mostly FS Middleware operating over multiple distributed nodes connected over a network.

## File Service Design

Distributed FS have many different kinds of design and implementation, each having its own characteristics and suited workloads.

A typical distributed file service model:



Typical caching techniques for distributed FS:

- Server-side caching:
  - Read-ahead
  - Delayed-write / Write-back v.s. Write-through v.s. Commit
- Client-side caching (Collective IO)

# Peer-to-Peer Systems (P2P)

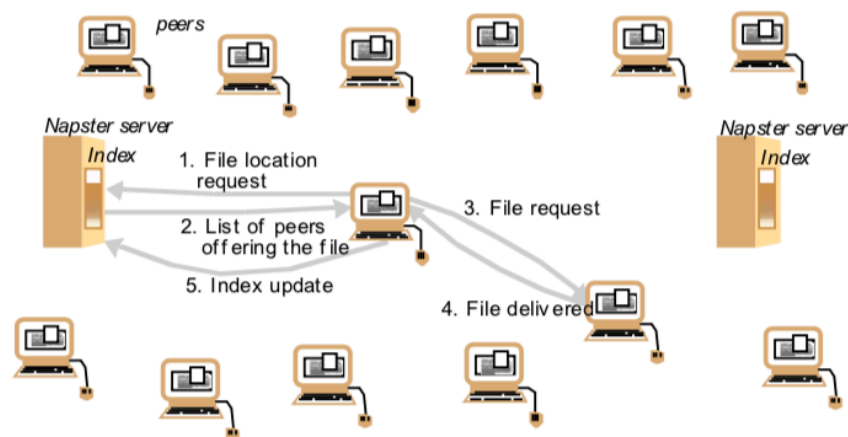
"Distributed Systems: Concepts & Design", Chapter 10

## Overview of P2P Systems

Purpose: Utilize data and computing resources available in personal computers on the Internet. Requires *Scalability, Reliability, and Security*.

Classification:

- Pure **P2P**: Peers act as equals; No central managing server; No central router
- Hybrid P2P: Has a central server that keeps info and responds requests



Key aspect of a P2P system is a set of algorithms for the placement and subsequent retrieval of information objects.

Examples of P2P systems:

- Music exchange services: Napster, ...
- File sharing: Freenet, Gnutella, Bit Torrent, ...
- *Overlay Networks* middleware running on top of the Internet (routing not necessarily specified by IP address): Pastry, Tapestry, ...

## Overlay Networks

Before overlay networks, P2P uses **Flooding**-Style Networks (using Time-To-Live limits, **TTL**):

1. Send query to all known neighbors;
2. Each neighbor checks to see whether they can reply (by matching keys);
3. If match, then reply; If not, ++hop\_count, and forward to its neighbors;
4. If hop\_count passed the TTL limit, stops forwarding.


[ **Distributed Hashing** ] **Consistent Hashing** and its classic implementation *Chord*: a *Ring*-based Overlay Network, nodes are linked in a circle, each having only two neighbors. A middleware layer takes the responsibility for routing request from any client to a host that holds the object to which the request is addressed (specified by a GUID). Nodes may:

- Join and leave the network
- Create / remove objects

A request on an overlay network is guaranteed to be handled.



A *Routing Table* can help reduce the number of hops of a message (by pushing a message to a node with longer common GUID prefix, instead of just searching in its leaf set, i.e., doing binary search). [READ HERE](#).



## Time & Global State

"Distributed Systems: Concepts & Design", Chapter 14

### Distributed Timing

Physical computer clocks (meaning the clock signal for electric circuits) are circuit oscillations at a well-defined frequency. Every designed number of oscillations trigger a timer interrupt. In a distributed system, timing is mostly *ambiguous*. Timing in different machines cannot guarantee to run at the same frequency.

Algorithms for keeping *absolute* times for distributed systems:

- Centralized timing for intranets:
  - **Christian's Algorithm:** For a cluster with one time server that has access to WWV time (UTC)
  - **Berkeley UNIX Algorithm:** For a cluster with no WWV access, time server polls other machines local times, computes an average, and tells all machines to adjust according to that
- Decentralized timing for intranets
- Internet time synchronization

Algorithms for keeping a *logical* time (not necessarily correct, but all machines agree):

- **Lamport Timestamps:** All machines agree on events' occurrence order, if they have message passing dependency
  - If A is the event of machine 1 sends a message, and B is the event of machine 2 receives it, we say A happens before B ( $A \rightarrow B$ )
  - If A and B happen in different machines and do not have dependency, we do not care about who happens first; we say they are concurrent ( $A \parallel B$ )
  - **Lamport's Clock Algorithm** ensures if  $A \rightarrow B$ , then all machines agree to  $C(A) < C(B)$ .  
[READ HERE](#) and refer to "Lec11.pdf".

### Global State

Global state of a cluster can be recorded by a **Snapshot**, reflecting a state in which the distributed system might have been.

- $H(i)$  is the History of node  $i$ , which is a vector of events  $\langle e_{i,0}, e_{i,1}, \dots \rangle$
- Global history  $H$  is the union of all individual histories  $H(0) \cup H(1) \cup \dots$
- A **Cut** of is a subset of the global history. A cut is *consistent* iff for every event it contains, it also contains all events happened before that event
  - Can contain a msg send event without the corresponding receive event
  - CANNOT contain a receive event without the corresponding send event

**Chandy-Lamport Algorithm:** for recording a consistent global state. [READ HERE](#) and refer to "Lec12.pdf".

## Coordination & Agreement

---

"Distributed Systems: Concepts & Design", Chapter 15

### Distributed Mutual Exclusion

Without shared variable, how can we achieve **Mutual Exclusion** only by message passing? Mutual exclusion requirements: 1. safety (mutual exclusiveness), 2. liveness (bounded waiting), 3.  $\rightarrow$  ordering.

Evaluating a distributed mutual exclusion algorithm:

- *Bandwidth*: number of messages sent in each `enter()` or `exit()` operation
- *Client delay*: waiting time at each `enter()` or `exit()`
- *Synchronization delay*: time between `exit()` and next process's `enter()`

Different algorithms to handle this:

1. **Central Server Algorithm**: use a central *token* server which queues entry requests, dequeue one request and gives the token to that client
  - Can violate " $\rightarrow$  ordering"
  - `enter()` takes 2 messages (request + receive token), `exit()` takes 1 message (return token)
  - Client delay depends on size of queue
  - Synchronization delay is 2 messages (return token + next one receives token)
2. **Ring-Based Algorithm**: peers arranged in a ring, enter when token is received and exit by passing on the token to next neighbor
  - Can violate " $\rightarrow$  ordering"
  - `enter()` has to wait for token to come, `exit()` takes 1 message (pass on token)
  - Client delay & Synchronization delay similar to entry
3. **Ricart-Agrawala Algorithm**: multicasting + Lamport timestamping. [READ HERE](#) and refer to "Lec13.pdf".

### Election Algorithms

How can we elect a unique process out of all distributed nodes? Election requirements: 1. safety (result either none or a process with largest identifier), 2. liveness (each process either picks the result or crashes).

Different algorithms to handle this:

1. **Ring-Based Election** (Chang-Roberts): [READ HERE](#) and refer to "Lec14.pdf".
2. **Bully Algorithm** (Garcia-Molina): can handle process failures. [READ HERE](#) and refer to "Lec14.pdf".

# Multicast Algorithms

**Multicast** is a sending scheme where we send a message to a group of nodes, and the message gets replicated only at path divergence.

I really don't understand the meaning and necessity of *reliable / ordered* multicast, so omitted here.

# Consensus Algorithms

How do a group of process agree to the same single value? That is called a **Consensus**. Consensus algorithms have to meet the following requirements:

- *Agreement*: Decided value of all correct processes is the same.
- *Integrity*: If all correct processes proposed the same value, then they must decide on that value.


BASE CASE - When processes cannot fail & communication cannot fail, the problem is easy: each processor broadcast its proposed value → waits until collected all  $N$  values → choose the majority one.

HARDER - When communication is reliable (no unintended message modifications) but processes can fail:

1. Two-Phase Commit (**2PC**), Three-Phase Commit (**3PC**): not 100% guaranteed, [READ HERE](#).
2. **Paxos** → Multi-Paxos → **Raft**: higher availability. [READ HERE](#) & my [blog](#).

EVEN HARDER - When message content can be wrong (modified / erroneous), i.e., **Byzantine Generals Problem**: Some of the processes can be faulty. If it is the commander (i.e., it is broadcasting), it may give out different values to different processes; If it is a general, it may relay a wrong value to others. (Message faults caused by unreliable communication links also count.)

- No solution when total number of processes  $N \leq 3f$ , where  $f$  is the number of faulty processes. A correct processor cannot tell who is faulty:



- For more information, read [Lamport's Paper](#) & [READ HERE](#), and refer to "Lec16.pdf".
- It is important for any modern real-time distributed system to be highly *Byzantine Fault Tolerant* (**BFT**); but, it's impossible at least for now to ensure 100% correctness, given that even the electric circuits can fail.

# Transactions Interface

"Distributed Systems: Concepts & Design", Chapter 16

# Definition of Transactions

A **Transaction** (事务) is a sequence of operations on a server that satisfy the **"ACID" rules** of databases:

- **Atomicity**: a transaction is either finished completely or not done at all; cannot end up in the middle (o.w., do *rollback*)

- For example, a failed file write may: a) write nothing at all, b) write a wrong value but checksums are used so that readers will detect that.
  - These are omission failures, and
  - The system must provide a failure model and corresponding recovery techniques
- But, this file write CANNOT write to the wrong block. That will be an arbitrary failure.
- **Consistency:** a transaction must bring the whole system from one valid state to another ("valid" is defined by a set of preset constraints, like some object's value must be  $\geq 0$ ); this is not the same term as in other contexts
- **Isolation:** allow concurrent transactions to act on the same object, w/o causing inatomicity / inconsistency
- **Durability:** whenever a transaction is done, its modification to objects are permanent and will not be lost even at system crashes, so that a crashed system can recover from the permanent storage

Providing a transaction *interface* to user applications can make the underneath concurrency control transparent to them. Users can simplify their application code logic by triggering transactions w/o worrying about violating the 4 rules.

Example on a banking account:

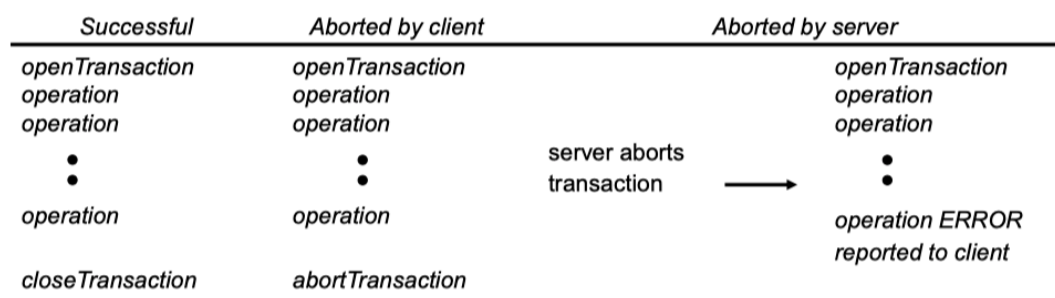
```

1  # Operations:
2  acc.withdraw(n)
3  acc.deposit(n)
4
5  # A Sample Transaction T:
6  def transaction_T(a, b, c):
7      a.withdraw(100)
8      b.deposit(100)    # A transfers $100 to B.
9      c.withdraw(200)
10     d.withdraw(200)   # C transfers $200 to D.

```

## Life History of a Transaction

The life history of a transaction must be one of the following situations:



- Operations are first done on volatile memory
- `close` = `commit`: flush the corresponding objects in volatile memory down to permanent storage
- `abort`: discard the changes in volatile memory

## Concurrency Control

First assumption: an **operation** itself must be atomic (called *synchronized operations*). Without this property, we cannot discuss the concurrency issues of transactions.

Problems that might happen w/o proper concurrency control on transactions:

1. **Lost update:** two transactions both read the old value

Transaction T:		Transaction U:	
<code>balance = b.getBalance();</code>		<code>balance = b.getBalance();</code>	
<code>b.setBalance(balance*1.1);</code>		<code>b.setBalance(balance*1.1);</code>	
<code>a.withdraw(balance/10)</code>		<code>c.withdraw(balance/10)</code>	
<code>balance = b.getBalance();</code>	\$200	<code>balance = b.getBalance();</code>	\$200
<code>b.setBalance(balance*1.1);</code>	\$220	<code>b.setBalance(balance*1.1);</code>	\$220
<code>a.withdraw(balance/10)</code>	\$80	<code>c.withdraw(balance/10)</code>	\$280

2. **Inconsistent retrieval:** a transaction reads an intermediate value of another transaction

Transaction V:		Transaction W:	
<code>a.withdraw(100)</code>		<code>aBranch.branchTotal()</code>	
<code>b.deposit(100)</code>			
<code>a.withdraw(100);</code>	\$100	<code>total = a.getBalance()</code>	\$100
		<code>total = total+b.getBalance()</code>	\$300
		<code>total = total+c.getBalance()</code>	
<code>b.deposit(100)</code>	\$300	:	

Correct concurrency of transactions should be **serial equivalent**: the result is the same as if they are processed serially. A *serially equivalent interleaving* is an interleaving of operations that makes the combined effect serial equivalent. Same effect here means:

- Read operations return the same value
- Object states are the same in the end

Operations that can conflict:

- `read(i)` and `write(i, x)` on the same object `i`
- `write(i, x)` and `write(i, y)` on the same object `i`

These operations MUST happen in the original order of transactions. Tools like DAG dependency graphs can help.

## Recoverability from Aborts

If a transaction *aborts*, the server must make sure that other transactions do not see any of its effects.

Problems that might happen w/o proper recoverability from aborts:

1. **Dirty read:** a transaction reads a modified value written by an aborted transaction

Transaction T:		Transaction U:	
<code>a.getBalance()</code>		<code>a.getBalance()</code>	
<code>a.setBalance(balance + 10)</code>		<code>a.setBalance(balance + 20)</code>	
<code>balance = a.getBalance()</code>	\$100	• U reads A's balance (which was set by T) and then commits	
<code>a.setBalance(balance + 10)</code>	\$110	<code>balance = a.getBalance()</code>	\$110
		<code>a.setBalance(balance + 20)</code>	\$130
T subsequently aborts.		<code>commit transaction</code>	
<code>abort transaction</code>			

2. **Premature write:** in a system that uses *before images*, it records the snapshot before a transaction and restore to the before image if that transaction aborts; if another transaction's write is done before the abortion, that write will get lost after restoration

Methods to ensure recoverability:

- *Delayed commit:* a commit is **delayed** until the commitment / abortion of other concurrent transactions are observed
  - If  $T$  aborts then  $U$  must also abort
  - May cause potential *cascading aborts*
- *Delayed read:* any read operation on  $i$  is delayed until other concurrent transactions who will apply writes on  $i$  are committed / aborted
  - This ensures only reading objects written by committed transactions, thus avoids cascading aborts
- *Delayed write:* any write operation on  $i$  is delayed until other concurrent transactions who will apply writes on  $i$  are committed / aborted
  - This prevents premature writes with before images
  - A system is said to implement **strict execution** if both read & write are delayed

## Nested Transactions

*Nested* Transactions are transactions composed of subtransactions recursively. To a *parent*, a subtransaction is atomic, while transactions at the same level can actually run concurrently just like over *flat* transactions.

Advantages:

- Additional parallelism
- More robust: a subtransaction failure can be properly handled by its parent, w/o restarting the whole transaction all over

Strict execution can be implemented using *strict two-phase locking*. Locks, *deadlocks* and *starvations* in transactions are omitted here. Already covered in details in the OS lectures.

**Distributed Transactions**, where a transaction accesses objects across multiple servers, are also omitted here. Refer to Chapter 17 of the book and "Lec{20,21}.pdf".

## Replications

"Distributed Systems: Concepts & Design", Chapter 18

**Replication** = Maintenance of copies of data at multiple computers. Can provide:

- Performance enhancement (rare)
- Fault tolerance
- High availability

## Replication Transparency

Replication of data should be **transparent** to users, which means that they appear as one single logical object, and different online users (disregard those disconnected copies) should see a *consistent* value.

Each *logical* object is implemented by a collection of *physial* copies, called **replicas**. All replicas on one device are managed by a *replica manager* (RM). The *frontend* (FE) talks to RMs and provides transparency to clients. Five phases in performing a request:

1. Issue request through FE
2. Coordination among RMs
  - Whether to apply / not
  - Its ordering relative to other requests (FIFO / Causal / Total-ordering)
3. Execution
4. Agreement: RMs reach a consensus on the effect of the request
5. Response back to FE

Often implemented using dynamic **Process Groups**.

## Fault-Tolerant Services

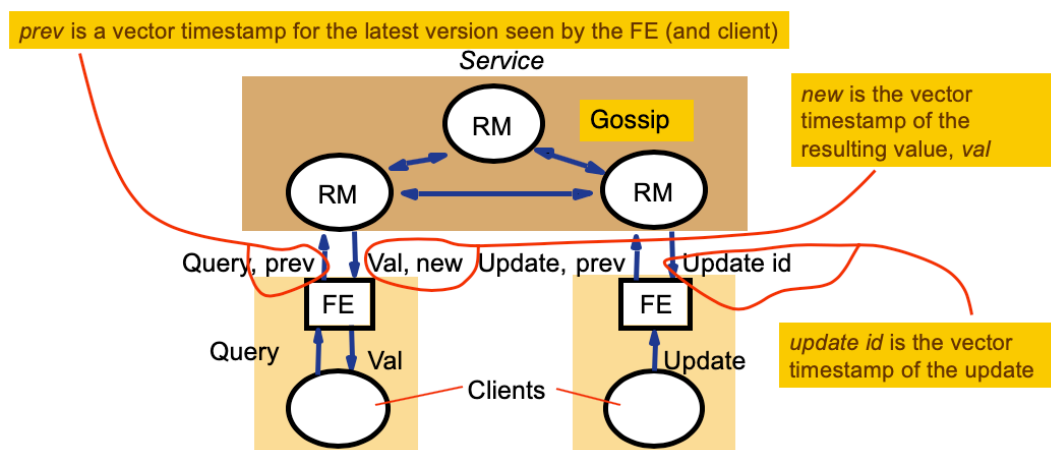
Replications enable **Fault-Tolerant Services**, where a service / data is still available even if up to  $f$  processes failed. Different models to do that:

- Passive (*primary-backup*) model: at any time there's a primary RM and others are backups
  - Needs  $f + 1$  RMs to achieve  $f$ -degree fault-tolerance
  - CANNOT handle Byzantine failures
- Active model: RMs all play the same role, and communicate with each other and compare the replicas it receives
  - Can use  $2f + 1$  RMs to handle up to  $f$  Byzantine failures
  - May have better performance
  - Needs a good consensus algorithm

## Highly Available Services

**High Availability** is a different goal from fault tolerance. We aim to give clients quick responses for as much of the time as possible, even if some results do not conform sequential consistency. (E.g., a disconnected client may accept some inconsistency and will fix that later.) Updates between RMs are propagated *lazily*, i.e., they have less agreement to enable shorter response time.

The **Gossip Architecture** is a way to provide high availability:



- Two types of operations: *Query* (read-only) and *Update* (will modify)

- FE sends operations to any chosen available RM, guaranteeing:
  - Each client gets consistent values *over time*, using vector timestamps
  - RMs eventually receive all updates, but they agree only by lazy gossips, so consistency is relaxed. Thus, a client may observe stale data at certain timepoint

## Security

---

"Distributed Systems: Concepts & Design", Chapter 11

Interesting names used in security protocols: Alice, Bob, Carol, Dave, Eve, Mallory, Sara ;)

This section should be found in Cryptology and Computer Security, so omitted here.