

Operating Systems

Author: Jose 胡冠洲 @ ShanghaiTech

Operating Systems

[Full-ver. Cheatsheet](#)

[Links](#)

Full-ver. Cheatsheet

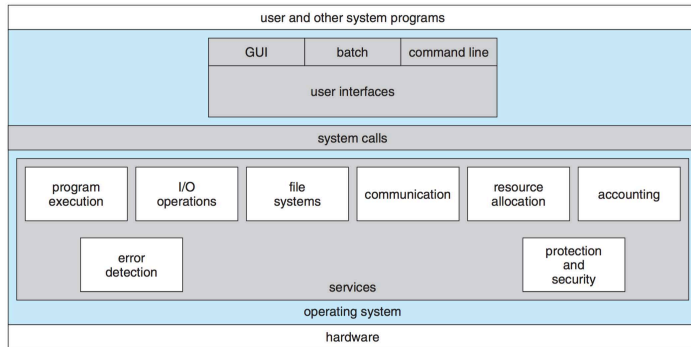
See below (page 2-6).

Links

- [Stanford CS140 Course Webpage](#)
- [Pintos Online Document](#)

What's an OS?

- | *Kernel* + (System programs) | User Apps
 - VM abstraction: APP-SW ↔ HW resources
 - Protection: SW | < CPU + Memory > | I/O
 - Loader for User programs
- 4 Fundamental OS Concepts
 - *Threads*
 - *Address Space*
 - *Processes*
 - *Dual Mode* Operation
- *bootstrap* program (*System Boot*):
 - Stored in firmware
 - Load kernel to run, after *SYSGEN*
- OS Services (provided via `syscall`)

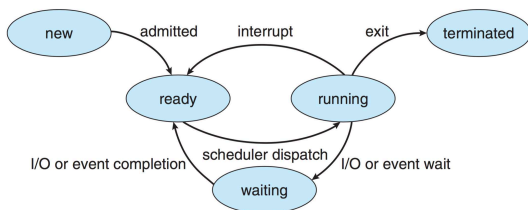


- Design Structures
 - Layered
 - Microkernel: Microkernel + system programs
 - Loadable Modules
 - Hybrid


Process

Instance (active) of an executing program (passive).

- + / –
 - + Protected from each other
 - + OS protected from them
 - – "Heavyweight", **different address spaces, page tables & file descriptors**
 - – Can only use kernel synchronization tools
- Process Control Block (**PCB**)
 - Status (new, running, ready, waiting, terminated)
 - Register state (when not ready)
 - Process ID (PID), User, Executable, Priority, ...
 - Execution time, ...
 - Memory space, Translation, ...



- Address space
 - ≠ base & bound
 - Owned by a process
 - Virtual, needs translator
- `fork()` : Child process is an **EXACT copy** of parent (separate address space).
 - Return value of `fork()` :
 - = $pid_{child} > 0$, then in parent
 - = 0, then in child
 - < 0, then error, in original
 - Waiting
 - *Zombie*: exit without parents currently waiting
 - *Orphan*: parent terminated without waiting
 - All processes are children of `init`
 - Have **Copy-on-write** technique



```
/* `syscall` APIs are as follows */
pid = fork();
exit(); // Terminate
pid = wait(&status); // Get status returned from child
abort(pid); // Terminate child process
exec("a_program"); // Flush the program being run currently
```

Threads

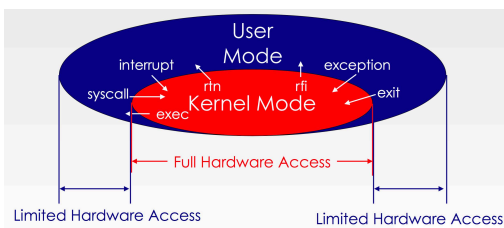
Single unique execution context.

- + / –
 - + Efficient, can use user synchronization
 - + "Lite Weight", **Share heap, static data & SAME code**
 - – Lacks protection
- Thread Control Block (**TCB**)
 - State (ready, running, blocked, terminated)
 - CPU register (when not ready)
 - Execution stack
- Internal / External events
 - Internal
 - Blocking on I/O
 - Waiting on other threads
 - Executes `yield()`
 - External
 - Interrupts (I/O, Timer)
- Thread library
 - User threads API
 - POSIX pthread
 - Windows
 - JAVA
 - Kernel threads: supported by kernel
- Thread model
 - Many → One
 - One → One
 - Many → Many
 - Two-Level model: M-M + allowing bound

```
pthread : POSIX thread
/* pthread APIs are as follows */
pthread_create(ind, NULL, (void *)worker, &tid);
pthread_join(&tid);
pthread_exit(); // Terminate current thread
pthread_kill(&tid); // Send a kill signal to
/* Mutex... */
/* Conditional Variables... */
```

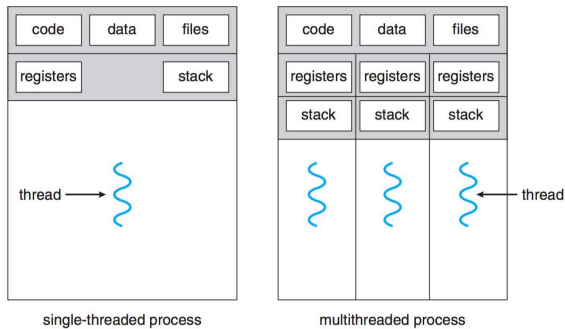
Dual mode & Context Switch

- *Context Switch*: Save & Load PCB / TCB
 - Have pure time overhead QAQ
- 2 Modes
 - **Kernel Mode**: Only "system" can access certain resources
 - **User mode**: User programs isolated from OS (and each other)
- Mode transfer
 - `syscall` : (e.g. `malloc()`)
 - `syscall` table, buried in run-time library
 - Process control
 - File manipulation
 - Device manipulation
 - Information maintenance
 - Communications
 - Protection
 - Parameter thransfering
 - Directly put in register
 - Memory Block and pass address to register
 - Via Stack: User → kernel stack
 - **Interrupts**: HW-invoked context switch (e.g. Timer)
 - Trap / Exception



Multi-xxx


- Definitions
 - Multiprocessing (core): Multiple CPUs (cores)
 - Multiprogramming: Multiple jobs or processes
 - Multithreading: Multiple threads per process
 - Time-sharing (Multitasking): switch frequently



- Concurrency**
 - Way: Can *multiplex in time*, virtual CPUs
 - Needs: Scheduler & Context Switch
- Parallelism**
 - Way: Data / Task parallelism
 - Needs: Multi-processors / Multicore / Hyperthreading / Pipelining

Interprocess Communication (IPC)

- Shared Memory
 - Unbounded Buffer
 - Bounded Buffer
- Direct Message Passing: Communication Link
 - Name each other explicitly
- Indirect Message Passing (*Mailboxes, ports*), possible solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive
 - Allow the system to select arbitrarily the receiver
 - Sender is notified who the receiver was
- Client-Server Communication
 - Sockets
 - Paired endpoints
 - Identified by IP address
 - Remote Procedure Call (*RPC*)
 - Message Passing between Clients & Server
 - Use Stubs to pass parameters
 - "Exactly Once"
 - Pipes
 - Anonymous (Ordinary)
 - Named (FIFO)



Synchronization

- Requirements
 - Mutual exclusion*
 - Making progress*
 - Bounded waiting*

```
/* Synch using *swap* */
void Initialize(int* lock) {
    *lock = 0;
}

void Acquire(int* lock) {
    int l = 1;
    do {
        swap(&l, lock);
    } while (l == 1);
}

void Release(int* lock) {
    *lock = 0;
}
```

- HW solutions
 - Atomic operations (i.e. unaltered Interrupts)
 - Test & Set, Page.210
 - Compare & Swap, Page
- SW solutions for
 - mutex : Busy waiting (*spinlock*) / *Blocking* lock
 - cond : Conditional variable
 - sema : Semaphore, Page.216

Specialty QAQ

- syscall parameters on thread stack, NOT kernel stack
 - Kernel pick a free interrupt number (e.g. 2)
 - fill function into IVT entry #2
 - User cause interrupt 2
 - Return from interrupt
- Even on single-processor, Multithreading speed up running
 - Since overlaps I/O with computation
- exec() only retains file descriptors
- Dual Mode + Virtual Address Translation = No over-writing ^-^
 - Since no privilege instructions allowed
- Process can wait / block on:
 - Acquiring a lock (sema, monitor)
 - Calling sleep
 - I/O call
 - wait() on child process
- Pintos Stuffs
 - Timer interrupts may preempt a thread
 - Idle thread: Simplicity for *scheduling* and *switch*
 - Interrupts off → Atomic operation
 - "magic" member must be at bottom of struct thread
 - User programs not implemented → ALL using kernel memory, crash all
- Sockets can be used either remotely or locally
- Interrupts can make locks
 - Single processor: ✓
 - Multiprocessors:** ✗
- fork() fails, wait() will immediately return
- Synchronization constraints
 - XXX must wait if XXX

```
/* Example of bounded-buffer solution */
1. Lock lock;
2. Condition dataready;
   Condition queueready;
3. Queue queue;

AddToQueue(item) {
    lock.Acquire();
    while (queue.isFull()) { /* WHILE LOOP!!! */
        queueready.wait(&lock);
    }
    queue.enqueue(item);
    dataready.signal();
    lock.Release();
}

RemoveFromQueue() {
    lock.Acquire(); // Get Lock
    while (queue.isEmpty()) {
        dataready.wait(&lock);
    }
    item = queue.dequeue();
    queueready.signal();
    lock.Release();
    return(item);
}
```

```
/* Barber Question */
void Barber () {
    while (true) {
        customerReady.P();
        accessWaitRoomSeats.P();
        numberOfFreeWaitRoomSeats += 1;
        accessWaitRoomSeats.V();
        cutHair();
        barberReady.V();
    }
}

void Customer () {
    accessWaitRoomSeats.P();
    if (numberOfFreeWaitRoomSeats > 0) {
        numberOfFreeWaitRoomSeats -= 1;
        accessWaitRoomSeats.V();
        customerReady.V();
        barberReady.P();
        getHairCut();
    } else {
        accessWaitRoomSeats.V();
        leaveWithoutHaircut();
    }
}
```

6 Synchronization

- **Critical Section Problem**
 - Solution **Requirements**: P194
 - *Software-based* Solutions:
 - **Peterson's**: P195
 - Bakery Algorithm
 - Providing Locks through **Hardware Atomic Instructions**:
 - TestAndSet() and Swap() : P197
 - Uniprocessor - Disable interrupts
 - Must have cache coherency
 - Providing **Semaphores** for Usage: P200
 - signal() and wait()
 - *Busy waiting* (spinlock) v.s. *Blocking*
- **Producer-Consumer Problem** (Bounded Buffer)
 - Shared data without synchronization solution
 - Allow at most $n - 1$ items in buffer
 - With counter *synchronized*: P205
- **Readers & Writers Problem**: P206
- **Dining Philosophers Problem**: P207
- Higher-level Synchronization - **Monitors**: P209
 - *Conditional* critical regions, ensure no deadlocks

7 Deadlocks

- **Definition**: P245
 - Deadlock \Rightarrow Starvation, \nLeftarrow
 - Starvation alone might end; Deadlocks cannot
- **Necessary Conditions**: P247
 - Mutual Exclusion
 - Hold + Wait
 - No Preemption
 - Circular Wait
- **System Model** (*Resource-Allocation Graph*): P249
- Methods of Handling Deadlocks: P252
 - **Deadlock Prevention**: P253
 - **Deadlock Avoidance**: P256
 - Banker's Algorithm
 - Ignore + Deadlock **Detection & Recovery**: P262
- Combined Approach
 - Hierarchical ordered resources classes
 - Use different technique for each class

8 Memory Management

- Background
 - **Base & Bound**: P277
 - **Address Binding**, *Logical* v.s. *Physical*: P278
 - Memory Management Unit (*MMU*)
 - Dynamic linking & loading: P280
- Primitive **Swapping**: P282
 - Pending I/O v.s. I/O to kernel space (Double Buffering)
 - Swap time = (Process size / Transfer rate) \times 2
- Primitive **Memory Allocation**: P284
 - External *Fragmentations*
- **Paging**: P288

- Internal *Fragmentations*
- Transition Look-aside Buffer (*TLB*)
- Effective Access Time =
 - $\text{Hit-ratio} \times (\text{TLB lookup time} + \text{Memory access time}) +$
 - $(1 - \text{Hit-ratio}) \times (\dots + \text{Page table access time})$
- *Multilevel* Paging & *Hashed* Paging: P299
- **Page Sharing** through Paging
 - Copy-on-Write (*CoW*), `vfork()` : P325
- **Segmentation**: P302

9 Virtual Memory

- **Demand Paging**: P319
 - Effective Access Time =
 - $(1 - \text{Page fault-ratio}) \times \text{Memory access time} +$
 - $\text{Page fault-ratio} \times (\text{Page fault overhead (almost 0)} +$
 - $\text{Swap time} + \text{Need swap out-ratio} \times \text{Swap time})$
 - Valid Bit
- **Page Replace** Algorithms, Modify Bit: P328
 - *FIFO*; *Belady's Anomaly*
 - Optimal
 - *LRU*
 - LRU Approximate
 - Reference (Access) Bit
 - *Second Chance (Clock)* Alg
 - Counting Based
 - LFU v.s. MFU
 - *Page Buffering*
- **Frame Allocation** Algorithms: P340
 - Equal v.s. Proportional
 - Global v.s. Local
- **Thrashing**: P343
 - Based on *Locality - Working Set & Page-fault Frequency*
- Other Issues: P357
 - Prepaging
 - Page size
 - TLB Reach
 - Program structure
 - I/O interlock

10-12 Storage Management

- **File Concepts, Operations, Types & Structures**: P373
 - Global open table v.s. Local open table; Offset is local!
- **Access Methods**: P382
 - *Sequential* v.s. *Direct*
- **Directory Concepts & Structures**: P385
 - Single level → Two level → Tree structure → Graph structure
- File **Protection**: P402
 - Access Control
 - *Consistency Semantics*
- **File System Layer Structure**: P411
- File System **Implementation**: P413
 - *Partitions & Mounting*
 - Virtual File System (*VFS*)
- **Disk Allocation** Methods: P421
 - Contiguous

- Linked + aFile-Allocation Table (*FAT*)
- *Indexed* (direct v.s. indirect)
- **Moving-head Disk:** P451
 - Transfer Rate
 - Random Access Time = Seek Time + Rotational Latency
 - Average Seek needs $\frac{1}{3}$ of overall Tracks
 - Average Latency needs $\frac{1}{2}$ of a circle
 - Average I/O time =
 - Random Access time +
 - (Amount to transfer / Transfer rate) +
 - Controller overhead
 - Effective Transfer Rate = Amount to transfer / Average I/O time
- **Disk Attachments:** P455
- **Disk Scheduling** Methods: P457
 - FCFS
 - SSTF
 - SCAN & C-SCAN
 - LOOK & C-LOOK
- **RAID** & Extensions: P470
 - Solaris ZFS system

Appendix

- Bakery Algorithm

```
int turn[n];
bool choosing[n];
int j;
while (1) {
    choosing[i] = true;
    turn[i] = 1 + max(turn[0], turn[1], ... turn[n-1]);
    choosing[i] = false;
    for (j = 0; j < n; j++) {
        if (j != i) {
            // Wait until thread j receives its number:
            while (choosing[j]);
            // Wait until all threads with smaller numbers
            // or with the same number but with higher priority
            // finish their work:
            while (turn[j] != 0 && ((turn[j], j) < (turn[i], i)));
            // (a, b) < (c, d) <=> (a < c) || (a = c && b < d)
        }
    }
    /* Critical Section. */
    turn[i] = 0;
    /* Remainder Section. */
}
```

- Shared data without synchronization Solution

```
/* Producer. */
while (true) {
    /* Produce an item `next_produced'. */
    while (((in + 1) % BUFFER_SIZE) == out);
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}

/* Consumer. */
while (true) {
    while ((in == out);
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* Consume the item `next_consumed'. */
}
```